

# Conception Objets Avancée

C++ threads

Giuseppe Lipari

CRISAL - Université de Lille

# Outline

Threads

Mutual exclusion

Condition variables

Asynchronous functions, future and promises

Performance issues

# Outline

## Threads

Mutual exclusion

Condition variables

Asynchronous functions, future and promises

Performance issues

# Creating a thread

- ▶ A thread is just an object of class `std::thread`
- ▶ To create a thread, you have to pass a **callable** (a function or a function object) to the constructor

---

```
void body() { ... }  
void g()  
{  
    std::thread some_thread(body);  
    ...  
    // wait for the thread to complete  
    some_thread.join();  
}
```

---

## Creating a thread, II

- ▶ The callable is copied into the thread object, so it is safe to destroy it afterwards

---

```
struct callable {  
    void operator()() { ... }  
};  
  
std::thread ok_function()  
{  
    callable x;  
    return std::thread(x);  
} // x is destroyed, but it has been copied, so this is ok
```

---

# Arguments

- Arguments are copied too, so (unless they are references), it is ok to destroy them afterwards

---

```
class A {...};  
  
void myfunction(A &param) { /* thread body */ }  
  
std::thread f()  
{  
    A obj;  
    return std::thread(myfunction, obj);  
}
```

---

- see [examples/thread\\_object\\_end.cpp](#)

# Examples

- ▶ Creating multiple threads
  - ▶ See [examples/threadex.cpp](#)
- ▶ A functional way to start threads
  - ▶ See [examples/thread\\_ret.cpp](#)
- ▶ A class that implements an active object
  - ▶ See [examples/thread\\_class.cpp](#)

# Copying threads

- ▶ Threads objects cannot be copied
  - ▶ However, they support the move constructor
  - ▶ In other words, threads are **never copied** (with the result of having two threads), but actually **moved**
  - ▶ this allows thread objects to be returned from functions without caring about ownership

---

```
std::thread make_thread()
{
    return std::thread(fun, arg);
}

int main() {
    std::thread th = make_thread();
    // there is only one thread active at this point
}
```

---



# Exceptions in threads

- ▶ If a callable passed to `std::thread` throws an exception, and the exception is not caught, the entire program is terminated
- ▶ See [examples/thread\\_exception.cpp](#)

# Destroying the thread object

- ▶ What happens if the object is destroyed before the thread has terminated?
  - ▶ see [examples/thread\\_nodetach.cpp](#)

---

```
void body()
{
    cout << "This is a thread" << endl;
    this_thread::sleep_for( chrono::seconds(1) );
    cout << "Am I still alive? " << endl;
}

void make_thread()
{
    thread mythread(body);
}

int main()
{
    make_thread();
    this_thread::sleep_for( chrono::seconds(3) );
    return 0;
}
```

---

# Detaching vs. joining

- ▶ In the previous example,
  - ▶ the **abort** is triggered by the fact that every thread must be joined by default
  - ▶ if we lose the reference to the thread object, we cannot join the thread anymore
- ▶ If this is the behaviour we need, then we have to **detach** the thread
  - ▶ this means that the thread continues execution, but it is not possible to join it anymore
  - ▶ it is possible to explicitly detach a thread by calling `detach` on the object.
  - ▶ in that case, the object becomes *not a thread*
  - ▶ if you do a join on a *not a thread* object, the join returns immediately
    - ▶ see how we can modify the above function to detach the thread

# Thread local

- ▶ Thread local variables are special variables that are allocated when a thread is created, and deleted when the thread terminates
  - ▶ There is exactly one copy per each thread

---

```
thread_local int index = 0;
void body(int x)
{
    index = x;
    cout << "My index is " << index << endl;
    ...
}
void fun()
{
    // the value printed depends on which
    // thread calls this function
    cout << index << endl;
}
```

---

# Outline

Threads

Mutual exclusion

Condition variables

Asynchronous functions, future and promises

Performance issues

# Mutexes

- ▶ C++ supports different kind of mutexes
  - ▶ **simple**, **recursive**, **shared**
  - ▶ each one can be timed or not
- ▶ Mutexes can be locked directly by using one of the following
  - ▶ `lock`, `try_lock`, `unlock`
  - ▶ A timed mutex also support `try_lock_for` and `try_lock_until` for a time-out
  - ▶ A shared mutex also supports `lock_shared` and `try_lock_shared`

# Simple mutex

- ▶ A mutex is used to avoid the possibility that several threads execute the same code
  - ▶ typically, this is sensible code that should not be executed at the same time by different threads, otherwise some inconsistency may arise
- ▶ An "stupid" example in the previous code(`examples/thread_local.cpp`): we do not want to mix the output of the different threads

# Guards

- ▶ Instead of accessing mutexes directly, it is useful to use guards for the RAIL technique
  - ▶ the simplest one is `lock_guard`

---

```
std::mutex m;  
  
void function()  
{  
    std::lock_guard<std::mutex> l(m);  
    // critical section of code  
    // if an exception is thrown,  
    // the mutex is automatically unlocked  
}
```

---



# Unique lock

- ▶ `unique_lock` allows also to directly release and acquire the ownership several times
  - ▶ when destroyed, it will unlock if necessary

---

```
std::mutex m;  
  
void function()  
{  
    std::lock_unique<std::mutex> l(m);  
    // critical section of code  
    l.unlock();  
    // normal code;  
    l.lock();  
    // critical section again  
}
```

---

# Recursive mutex

- ▶ A recursive mutex can be useful if we need to lock the same mutex several times from the thread
  - ▶ for example in a recursive function

---

```
int recursive_function(int x)
{
    if (x == 0) {
        std::unique_lock<std::recursive_mutex> l(m);
        std::cout << "End of recursion" << std::endl;
        return x;
    }
    else {
        std::unique_lock<std::recursive_mutex> l(m);
        cout << "This is recursive_function(" << x << ");" << endl;
        global++;
        return recursive_function(--x);
    }
}
```

---

## Another example

- ▶ in the following example, function `f()` can be called from the `main()` or from function `g()`

---

```
class Shared {
private:
    std::recursive_mutex m;
public:
    Shared() {}
    Shared(const Shared &) = delete;
    Shared & operator=(const Shared &) = delete;

    void f() {
        std::unique_lock<std::recursive_mutex> l(m);
        std::cout << "Inside function f() " << std::endl;
    }
    void g() {
        std::unique_lock<std::recursive_mutex> l(m);
        std::cout << "Inside function g()" << std::endl;
        f();
    }
};
```

---

# Direct locking

- ▶ It is possible to lock/unlock mutexes using global functions
- ▶ This can be used when we have to lock several mutexes at once

---

```
struct bank_account {  
    explicit bank_account(int balance) : balance(balance) {}  
    int balance;  
    std::mutex m;  
};  
  
void transfer(bank_account &from, bank_account &to, int amount)  
{  
    // attempts to lock both mutexes without deadlock  
    std::lock(from.m, to.m);  
    // code  
    std::unlock(from.m, to.m);  
}
```

---

## Executing code only once

- It is possible to specify that a certain function must be called only once

---

```
std::once_flag flag;

void do_once()
{
    std::call_once(flag, [](){ std::cout << "Called once" << std::endl; });
}

int main()
{
    std::thread t1(do_once);
    std::thread t2(do_once);
    std::thread t3(do_once);

    t1.join();
    t2.join();
    t3.join();
}
```

---

# Shared mutex

- ▶ A shared mutex allows to implement the reader/writer paradigm
- ▶ Suppose you have several threads that want to read from a data structure and a few writers that want to write on the data structure
  - ▶ we must avoid that two writers, or a reader and a writer access the data structure at the same time
  - ▶ however, readers can actually read from the data structure concurrently

	Reader	Writer
Reader	Ok	NO
Writer	NO	NO

# Shared mutex - II

- ▶ Using simple mutexes
  - ▶ every thread is blocked if a thread is access a data structure
  - ▶ this means that two readers cannot read at the same time
  - ▶ This is a safe solution, but it is not very efficient
- ▶ Shared mutex
  - ▶ writers use the normal lock operations, while readers use the special `shared_lock` operation
  - ▶ a `shared_lock` operation is blocking only if someone has already done a normal lock operation
- ▶ Only available from C++17

# Outline

Threads

Mutual exclusion

Condition variables

Asynchronous functions, future and promises

Performance issues



# Conditions

- ▶ Condition variables are used to synchronize threads
  - ▶ They must always be paired with one mutex
- ▶ Interface:
  - ▶ default constructor only, cannot be copied or moved
  - ▶ wait: takes a `unique_lock` as parameter, and unlocks it while the thread is blocked

---

```
std::mutex m;  
std::condition_variable cv;  
  
int fun()  
{  
    std::unique_lock l(m);  
    while (condition) cv.wait(l);  
  
    // can execute critical section code  
}
```

---

# Interface

- ▶ Interface (cont)
  - ▶ `wait_for`: a timed version of the wait
  - ▶ `wait_until`: another timed version
  - ▶ `notify`: unblocks the first thread in the queue
  - ▶ `notify_all`: unblocks all threads in the queue
  - ▶ `wait(uunlock, pred)` is a shortcut for `while (pred) wait(uunlock);`
- ▶ The while is necessary because of **spurious wake-ups**
  - ▶ It may happen that the thread is unblocked, but the condition is false

# The mailbox example

- ▶ The following is a didactic example to demonstrate all we have seen until now
- ▶ See `examples/mailbox.hpp` and `examples/thread_mailbox.cpp`

# Synchronization barrier

- ▶ Let us write a synchronization barrier
  - ▶ threads can continue execution only when all thread reach the barrier
- ▶ See [examples/thread\\_group.cpp](#)

# Outline

Threads

Mutual exclusion

Condition variables

Asynchronous functions, future and promises

Performance issues

# Asynchronous calls

- ▶ In C++ it is possible to defer the call to a function to another thread, or later on when the result is actually needed (lazy evaluation).
  - ▶ this is done through the `async` function

---

```
template< class Function, class... Args>
std::future<typename std::result_of<Function(Args...)>::type>
    async( Function&& f, Args&&... args );
```

```
template< class Function, class... Args >
std::future<typename std::result_of<Function(Args...)>::type>
    async( std::launch policy, Function&& f, Args&&... args );
```

---

- ▶ Launch policy can be
  - ▶ `std::launch::async`: a thread is created
  - ▶ `std::launch::deferred`: lazy evaluation
  - ▶ If both are specified in or, the actual behaviour is implementation dependent

# Asynchronous calls and futures

- ▶ An async call returns a `future`, that is an object that at some point will contain the result of the function
- ▶ If a thread tries to obtain the value before the function has completed, it blocks
  - ▶ see `examples/parallel_sum.cpp`

# Async and futures

- ▶ **Warning:** if the future returned from `async` is not copied into a variable, the destructor waits for the function to terminate
- ▶ In other words, it does not parallelise anything

---

```
std::async(std::launch::async, []{ return f(); });  
// temporary's dtor waits for f()  
std::async(std::launch::async, []{ return g(); });  
// does not start until f() completes
```

---



# Future validity

- ▶ When you call `get` on a future, you are actually **moving out** the result from the future
- ▶ After calling `get`, the future becomes *invalid*
- ▶ To check the validity of a future, you must call `valid`
- ▶ If you call `get` on an invalid future, the behaviour is **undefined**
  - ▶ most implementation will throw a `std::future_error` exception

# Packaged task

- ▶ A packaged task is similar to an `async`,
- ▶ it is an object that wraps a function, that can be later executed in several ways:
  - ▶ as a thread
  - ▶ as a normal function
- ▶ we can obtain the corresponding future with `get_future`

## Example of packaged task into a thread

- ▶ See [http://en.cppreference.com/w/cpp/thread/packaged\\_task](http://en.cppreference.com/w/cpp/thread/packaged_task)

# Outline

Threads

Mutual exclusion

Condition variables

Asynchronous functions, future and promises

Performance issues

# Parallel code

- ▶ It is quite difficult to write fast parallel code, because of the many sharing parts
- ▶ Let's try to write a very simple parallel program, and let's see if we can make it fast
- ▶ We want to compute how many elements in a matrix are odd.
- ▶ Sequential code:

---

```
int odds = 0;
for( int i = 0; i < DIM; ++i )
    for( int j = 0; j < DIM; ++j )
        if( matrix[i*DIM + j] % 2 != 0 )
            ++odds;
```

---

# Naive parallel code

---

```
void body(unsigned start, unsigned end, int &result) {
    for (unsigned i=start; i<end; ++i) {
        for (unsigned j=0; j<NELEM; ++j)
            if (matrix[i][j] % 2) ++result;
    }
}

...
vector<int> result(nthreads);
for (unsigned p = 0; p<nthreads; ++p) {
    ...
    th.emplace_back(thread(body, start, end, result[p]));
}
total = accumulate(result.begin(), result.end(), 0);
```

---

# Analysis


- ▶ If all threads write on global variables there can be a lot of load on the cache.
- ▶ It may be better to operate on local variables, and then write the result only at the end

---

```
void body(unsigned start, unsigned end, int &result)
{
    int temp = 0;
    for (unsigned i=start; i<end; ++i)
        for (unsigned j=0; j<NELEM; ++j)
            if (matrix[i][j] % 2) ++temp;
    result = temp;
}
```

---

# Comparison





# Conclusions

- ▶ Effective parallelization is not so easy
  - ▶ the structure of the actual hw architecture is very important
  - ▶ L1 and L2 cache size
  - ▶ number of cores, number of hyperthreads
  - ▶ branch prediction
  - ▶ cache pre-load algorithms
  - ▶ size of the data structures (do they fit in the cache?)
  - ▶ compiler
  - ▶ etc.
- ▶ Only approach: try different solutions and and measure!